

SQL guide: Getting started

Just as humans use different languages to communicate with others, so do computers. **Structured Query Language** (or **SQL**, often pronounced “sequel”) enables data analysts to talk to their databases. SQL is one of the most useful data analyst tools, especially when working with large datasets in tables. It can help you investigate huge databases, track down text (referred to as strings) and numbers, and filter for the exact kind of data you need—much faster than a spreadsheet can.

If you haven’t used SQL before, this reading will help you learn the basics so you can appreciate how useful SQL is and how useful SQL queries are in particular. You will be writing SQL queries in no time at all.

What is a query?

A **query** is a request for data or information from a database. When you query databases, you use SQL to communicate your question or request. You and the database can always exchange information as long as you speak the same language.

Every programming language, including SQL, follows a unique set of guidelines known as **syntax**. Syntax is the predetermined structure of a language that includes all required words, symbols, and punctuation, as well as their proper placement. As soon as you enter your search criteria using the correct syntax, the query starts working to pull the data you’ve requested from the target database. The syntax of every SQL query is the same:

- Use **SELECT** to choose the columns you want to return.
- Use **FROM** to choose the tables where the columns you want are located.
- Use **WHERE** to filter for certain information.

A SQL query is like filling in a template. You will find that if you are writing a SQL query from scratch, it is helpful to start a query by writing the **SELECT**, **FROM**, and **WHERE** keywords in the following format:

SELECT

FROM

WHERE

Next, enter the table name after the **FROM**; the table columns you want after the **SELECT**; and, finally, the conditions you want to place on your query after the **WHERE**. Make sure to add a new line and indent when adding these, as shown below:

SELECT	Specifies the columns from which to retrieve data
FROM	Specifies the table from which to retrieve data
WHERE	Specifies criteria that the data must meet

Following this method each time makes it easier to write SQL queries. It can also help you make fewer syntax errors.

Example of a query

Here is how a simple query would appear in BigQuery, a data warehouse on the Google Cloud Platform.

```

1 SELECT first_name
2 FROM customer_data.customer_name
3 WHERE first_name = 'Tony'

```

The above query uses three commands to locate customers with the `first_name`, 'Tony':

1. **SELECT** the column named `first_name`

FROM a table named `customer_name` (in a dataset named `customer_data`)

2. (The dataset name is always followed by a dot, and then the table name.)
3. But only return the data **WHERE** the `first_name` is 'Tony'

The results from the query might be similar to the following:

first_name
Tony
Tony
Tony

As you can conclude, this query had the correct syntax, but wasn't very useful after the data was returned.

Multiple columns in a query

Of course, as a data professional, you will need to work with more data beyond customers named Tony. Multiple columns that are chosen by the same **SELECT** command can be indented and grouped together.

If you are requesting multiple data fields from a table, you need to include these columns in your **SELECT** command. Each column is separated by a comma as shown below:

```

1 SELECT
2   ColumnA,
3   ColumnB,
4   ColumnC
5 FROM
6   Table where the data lives
7 WHERE
8   Certain condition is met

```

Here is an example of how it would appear in BigQuery:

```

1  SELECT
2      customer_id,
3      first_name,
4      last_name
5  FROM
6      customer_data.customer_name
7  WHERE
8      first_name = 'Tony'

```

The above query uses three commands to locate customers with the `first_name`, 'Tony'.

1. `SELECT` the columns named `customer_id`, `first_name`, and `last_name`
- `FROM` a table named `customer_name` (in a dataset named `customer_data`)
2. (The dataset name is always followed by a dot, and then the table name.)
3. But only return the data `WHERE` the `first_name` is 'Tony'

The only difference between this query and the previous one is that more data columns are selected. The previous query selected `first_name` only while this query selects `customer_id` and `last_name` in addition to `first_name`. In general, it is a more efficient use of resources to select only the columns that you need. For example, it makes sense to select more columns if you will actually use the additional fields in your `WHERE` clause. If you have multiple conditions in your `WHERE` clause, they may be written like this:

```

1  SELECT
2      ColumnA,
3      ColumnB,
4      ColumnC
5  FROM
6      Table where the data lives
7  WHERE
8      Condition 1
9      AND Condition 2
10     AND Condition 3

```

Notice that unlike the `SELECT` command that uses a comma to separate fields / variables / parameters, the `WHERE` command uses the `AND` statement to connect conditions. As you become a more advanced writer of queries, you will make use of other connectors / operators such as `OR` and `NOT`.

Here is a BigQuery example with multiple fields used in a `WHERE` clause:

```

1  SELECT
2      customer_id,
3      first_name,
4      last_name
5  FROM
6      customer_data.customer_name
7  WHERE
8      customer_id > 0
9      AND first_name = 'Tony'
10     AND last_name = 'Magnolia'

```

The above query uses three commands to locate customers with a valid (greater than 0), `customer_id` whose `first_name` is 'Tony' and `last_name` is 'Magnolia'.

1. **SELECT** the columns named `customer_id`, `first_name`, and `last_name`

FROM a table named `customer_name` (in a dataset named `customer_data`)

2. (The dataset name is always followed by a dot, and then the table name.)

3. But only return the data **WHERE** `customer_id` is greater than 0, `first_name` is Tony, and `last_name` is Magnolia.

Note that one of the conditions is a logical condition that checks to see if `customer_id` is greater than zero.

If only one customer is named Tony Magnolia, the results from the query could be:

customer_id	first_name	last_name
1967	Tony	Magnolia

If more than one customer has the same name, the results from the query could be:

customer_id	first_name	last_name
1967	Tony	Magnolia
7689	Tony	Magnolia

Key takeaways

The **SELECT**, **FROM**, and **WHERE** clauses are the essential building blocks of SQL queries. Queries with multiple fields will become simpler after you practice writing your own SQL queries later in the program.

Endless SQL possibilities

You have learned that a SQL query uses **SELECT**, **FROM**, and **WHERE** to specify the data to be returned from the query. This reading provides more detailed information about formatting queries, using **WHERE** conditions, selecting all columns in a table, adding comments, and using aliases. All of these make it easier for you to understand (and write) queries to put SQL in action. The last section of this reading provides an example of what a data analyst would do to pull employee data for a project.

Capitalization, indentation, and semicolons

You can write your SQL queries in all lowercase and don't have to worry about extra spaces between words. However, using capitalization and indentation can help you read the information more easily. Keep your queries neat, and they will be easier to review or troubleshoot if you need to check them later on.

```
1  SELECT field1
2  FROM table
3  WHERE field1 = condition;
```

Notice that the SQL statement shown above has a semicolon at the end. The semicolon is a statement terminator and is part of the American National Standards Institute (ANSI) SQL-92 standard, which is a recommended common syntax for adoption by all SQL databases. However, not all SQL databases have adopted or enforce the semicolon, so it's possible you may come across some SQL statements that aren't terminated with a semicolon. If a statement works without a semicolon, it's fine.

WHERE conditions

In the query shown above, the **SELECT** clause identifies the column you want to pull data from by name, `field1`, and the **FROM** clause identifies the `table` where the column is located by name, `table`. Finally, the **WHERE** clause narrows your query so that the database returns only the data with an exact value match or the data that matches a certain condition that you want to satisfy.

For example, if you are looking for a specific customer with the last name Chavez, the **WHERE** clause would be:

```
WHERE field1 = 'Chavez'
```

However, if you are looking for all customers with a last name that begins with the letters "Ch," the **WHERE** clause would be:

```
WHERE field1 LIKE 'Ch%'
```

You can conclude that the **LIKE** clause is very powerful because it allows you to tell the database to look for a certain pattern! The percent sign `%` is used as a wildcard to match

one or more characters. In the example above, both **Chavez** and **Chen** would be returned. Note that in some databases an asterisk ***** is used as the wildcard instead of a percent sign **%**.

SELECT all columns

Can you use `SELECT *` ?

In the example, if you replace `SELECT field1` with `SELECT *`, you would be selecting all of the columns in the table instead of the `field1` column only. From a syntax point of view, it is a correct SQL statement, but you should use the asterisk ***** sparingly and with caution. Depending on how many columns a table has, you could be selecting a tremendous amount of data. Selecting too much data can cause a query to run slowly.

Comments

Some tables aren't designed with descriptive enough naming conventions. In the example, `field1` was the column for a customer's last name, but you wouldn't know it by the name. A better name would have been something such as `last_name`. In these cases, you can place comments alongside your SQL to help you remember what the name represents. Comments are text placed between certain characters, `/*` and `*/`, or after two dashes `--`) as shown below.

```
1  SELECT
2  |    field1 /* this is the last name column */
3  FROM
4  |    table -- this is the customer data table
5  WHERE
6  |    field1 LIKE 'ch%';
```

Comments can also be added outside of a statement as well as within a statement. You can use this flexibility to provide an overall description of what you are going to do, step-by-step notes about how you achieve it, and why you set different parameters/conditions.

```
1  -- This is an important query used later to join with the accounts table
2  SELECT
3  |    rowkey, -- key used to join with account_id
4  Info.date, -- date is in string format YYYY-MM-DD HH:MM:SS
5  Info.code -- e.g., 'pub-###'
6
7  FROM Publishers
```

The more comfortable you get with SQL, the easier it will be to read and understand queries at a glance. Still, it never hurts to have comments in a query to remind yourself of what you're trying to do. This also makes it easier for others to understand your query if your query is shared. As your queries become more and more complex, this practice will save you a lot of time and energy to understand complex queries you wrote months or years ago.

Example of a query with comments

Here is an example of how comments could be written in BigQuery:

```
1  -- Pull basic information from the customer table
2  SELECT
3      customer_id, --main ID used to join with customer_addresses
4      first_name, --customer's first name from loyalty program
5      last_name --customer's last name
6  FROM
7      customer_data.customer_name
```

In the above example, a comment has been added before the SQL statement to explain what the query does. Additionally, a comment has been added next to each of the column names to describe the column and its use. Two dashes -- are generally supported. So it is best to use -- and be consistent with it. You can use # in place of -- in the above query, but # is not recognized in all SQL versions; for example, MySQL doesn't recognize #. You can also place comments between /* and */ if the database you are using supports it.

As you develop your skills professionally, depending on the SQL database you use, you can pick the appropriate comment delimiting symbols you prefer and stick with those as a consistent style. As your queries become more and more complex, the practice of adding helpful comments will save you a lot of time and energy to understand queries that you may have written months or years prior.

Aliases

You can also make it easier on yourself by assigning a new name or **alias** to the column or table names to make them easier to work with (and avoid the need for comments). This is done with a SQL AS clause. In the example below, aliases are used for both a table name and a column. Within the database, the table is called `actual_table_name` and the column in that table is called `actual_column_name`. They are aliased as `my_table_alias` and `my_column_alias`, respectively. These aliases are good for the duration of the query only. An alias doesn't change the actual name of a column or table in the database.

Example of a query with aliases

```
1 SELECT
2 | my_table_alias.actual_column_name AS my_column_alias
3 FROM
4 | actual_table_name AS my_table_alias
```

Putting SQL to work as a data analyst

Imagine you are a data analyst for a small business and your manager asks you for some employee data. You decide to write a query with SQL to get what you need from the database.

You want to pull all the columns: **empID**, **firstName**, **lastName**, **jobCode**, and **salary**. Because you know the database isn't that big, instead of entering each column name in the **SELECT** clause, you use **SELECT ***. This will select all the columns from the **Employee** table in the **FROM** clause.

```
1 SELECT
2 | *
3 FROM
4 | Employee
```

Now, you can get more specific about the data you want from the **Employee** table. If you want all the data about employees working in the 'SFI' job code, you can use a **WHERE** clause to filter out the data based on this additional requirement.

Here, you use:

```
1 SELECT
2 | *
3 FROM
4 | Employee
5 WHERE
6 | jobCode = 'SFI'
```

A portion of the resulting data returned from the SQL query might look like this:

empID	firstName	lastName	jobCode	salary
0002	Homer	Simpson	SFI	15000
0003	Marge	Simpson	SFI	30000
0034	Bart	Simpson	SFI	25000
0067	Lisa	Simpson	SFI	38000
0088	Ned	Flanders	SFI	42000
0076	Barney	Gumble	SFI	32000

Suppose you notice a large salary range for the 'SFI' job code. You might like to flag all employees in all departments with lower salaries for your manager. Because interns are also included in the table and they have salaries less than \$30,000, you want to make sure your results give you only the full time employees with salaries that are \$30,000 or less. In other words, you want to exclude interns with the 'INT' job code who also earn less than \$30,000. The **AND** clause enables you to test for both conditions.

You create a SQL query similar to below, where **<>** means "does not equal":

```

1  SELECT
2  |   *
3  FROM
4  |   Employee
5  WHERE
6  |   jobCode <> 'INT'
7  |   AND salary <= 30000;

```

The resulting data from the SQL query might look like the following (interns with the job code **INT** aren't returned):

empID	firstName	lastName	jobCode	salary
0002	Homer	Simpson	SFI	15000
0003	Marge	Simpson	SFI	30000
0034	Bart	Simpson	SFI	25000
0108	Edna	Krabappel	TUL	18000
0099	Moe	Szyslak	ANA	28000

With quick access to this kind of data using SQL, you can provide your manager with tons of different insights about employee data, including whether employee salaries across the business are equitable. Fortunately, the query shows only an additional two employees might need a salary adjustment and you share the results with your manager.

Pulling the data, analyzing it, and implementing a solution might ultimately help improve employee satisfaction and loyalty. That makes SQL a pretty powerful tool.

Resources to learn more

Nonsubscribers may access these resources for free, but if a site limits the number of free articles per month and you already reached your limit, bookmark the resource and come back to it later.

- [W3Schools SQL Tutorial](#): If you would like to explore a detailed tutorial of SQL, this is the perfect place to start. This tutorial includes interactive examples you can edit, test, and recreate. Use it as a reference or complete the whole tutorial to practice using SQL. Click the green **Start learning SQL now** button or the **Next** button to begin the tutorial.
- [SQL Cheat Sheet](#): For more advanced learners, go through this handy 3-page resource to gain an overview of additional SQL functions and formulas. By the time you are finished looking through the cheat sheet, you will know a lot more about the various SQL techniques and will be prepared to use it for business analysis and other tasks.

Key takeaways

SQL queries use **SELECT**, **FROM**, and **WHERE** to specify the data to be returned from the query. Capitalization, indentation, and semicolons are useful for making your SQL queries easier to read. In addition, comments can be added to explain queries to others. As you progress through this course, you will continue discovering many ways in which SQL can be a very powerful tool for retrieving, analyzing, and interpreting data.

Plan a data visualization

Earlier, you learned that **data visualization** is the graphical representation of information. As a data analyst, you will want to create visualizations that make your data easy to understand and interesting to look at. Because of the importance of data visualization, most data analytics tools (such as spreadsheets and databases) have a built-in visualization component while others (such as Tableau) specialize in visualization as their primary value-add. In this reading, you will explore the steps involved in the data visualization process and a few of the most common data visualization tools available.



Steps to plan a data visualization

Let's go through an example of a real-life situation where a data analyst might need to create a data visualization to share with stakeholders. Imagine you're a data analyst for a clothing distributor. The company helps small clothing stores manage their inventory, and sales are booming. One day, you learn that your company is getting ready to make a major update to its website. To guide decisions for the website update, you're asked to analyze data from the existing website and sales records. Let's go through the steps you might follow.

Step 1: Explore the data for patterns

First, you ask your manager or the data owner for access to the current sales records and website analytics reports. This includes information about how customers behave on the company's existing website, basic information about who visited, who bought from the company, and how much they bought.

While reviewing the data you notice a pattern among those who visit the company's website most frequently: geography and larger amounts spent on purchases. With further analysis, this information might explain why sales are so strong right now in the northeast—and help your company find ways to make them even stronger through the new website.

Step 2: Plan your visuals

Next it is time to refine the data and present the results of your analysis. Right now, you have a lot of data spread across several different tables, which isn't an ideal way to share your results with management and the marketing team. You will want to create a data visualization that explains your

findings quickly and effectively to your target audience. Since you know your audience is sales oriented, you already know that the data visualization you use should:


- Show sales numbers over time
- Connect sales to location
- Show the relationship between sales and website use
- Show which customers fuel growth

Step 3: Create your visuals

Now that you have decided what kind of information and insights you want to display, it is time to start creating the actual visualizations. Keep in mind that creating the right visualization for a presentation or to share with stakeholders is a process. It involves trying different visualization formats and making adjustments until you get what you are looking for. In this case, a mix of different visuals will best communicate your findings and turn your analysis into the most compelling story for stakeholders. So, you can use the built-in chart capabilities in your spreadsheets to organize the data and create your visuals.




Line charts can track sales over time



Maps can connect sales to locations



Donut charts can show customer segments



Bar charts can compare total visitors and visitors that make a purchase

1) line charts can track sales over time

- 2) maps can connect sales to locations
- 3) donut charts can show customer segments
- 4) bar charts can compare total visitors that make a purchase

Build your data visualization toolkit

There are many different tools you can use for data visualization.

- You can use the visualizations tools in your spreadsheet to create simple visualizations such as line and bar charts.
- You can use more advanced tools such as Tableau that allow you to integrate data into dashboard-style visualizations.
- If you're working with the programming language R you can use the visualization tools in RStudio.

Your choice of visualization will be driven by a variety of drivers including the size of your data, the process you used for analyzing your data (spreadsheet, or databases/queries, or programming languages). For now, just consider the basics.


Spreadsheets (Microsoft Excel or Google Sheets)

In our example, the built-in charts and graphs in spreadsheets made the process of creating visuals quick and easy. Spreadsheets are great for creating simple visualizations like bar graphs and pie charts, and even provide some advanced visualizations like maps, and waterfall and funnel diagrams (shown in the following figures).

But sometimes you need a more powerful tool to truly bring your data to life. Tableau and RStudio are two examples of widely used platforms that can help you plan, create, and present effective and compelling data visualizations.

Visualization software (Tableau)


Tableau is a popular data visualization tool that lets you pull data from nearly any system and turn it into compelling visuals or actionable insights. The platform offers built-in visual best practices, which makes analyzing and sharing data fast, easy, and (most importantly) useful. Tableau works well with a wide variety of data and includes an interactive dashboard that lets you and your stakeholders click to explore the data interactively.



You can start exploring Tableau from the [How-to Video](#) resources. Tableau Public is free, easy to use, and full of helpful information. The Resources page is a one-stop-shop for how-to videos, examples, and datasets for you to practice with. To explore what other data analysts are sharing on Tableau, visit the [Viz of the Day](#) page where you will find beautiful visuals ranging from an overview of the [Lighthouses of Greece](#) to [Who's Talking in Popular Films](#).

Programming language (R with RStudio)

A lot of data analysts work with a programming language called R. Most people who work with R end up also using RStudio, an integrated developer environment (IDE), for their data visualization needs. As with Tableau, you can create dashboard-style data visualizations using RStudio.



Check out their website to learn more about [RStudio](#).

You could easily spend days exploring all the resources provided at RStudio.com, but the [RStudio Cheatsheets](#) and the [RStudio Visualize Data Primer](#) are great places to start. When you have more time, check out the webinars and videos which offer advice and helpful perspectives for both beginners and advanced users.

Key takeaways

The best data analysts use lots of different tools and methods to visualize and share their data. As you continue learning more about data visualization throughout this course, be sure to stay curious, research different options, and continuously test new programs and platforms to help you make the most of your data.